


## Critical Section in a Process: Environment

- $n$  processes all competing to use some shared variables (data).
- Each process has multiple code segments, each called a **critical section**, in which the shared data is accessed/modified.
- Problem – Make sure that the three critical section problem properties are always satisfied.

1

## Program Structure (with a Single Critical Section) (simplification for the time being)



General structure of process  $P_i$  (*only in this part of slides*)

2

## Two-Process ( $P_0$ and $P_1$ ) Synchronization: Algorithm 1

```

int turn;          /* turn is a shared variable; initially, turn = 0
                    turn == i  $\Rightarrow$   $P_i$  can enter its critical section */

 $P_i$  : repeat
    while (turn != i) do no-op;    /* turn=0  $\Rightarrow$   $P_0$  can enter CS
    critical section                  turn=1  $\Rightarrow$   $P_1$  can enter CS */

    turn := j;                      /* j=1 - i, 0  $\leq$  i, j  $\leq$  1 */
    remainder section

until false;

```

Correct? Satisfaction of (1) mutual exclusion?  
 (2) Progress?  
 (3) Bounded Waiting?

3

## Two-Process ( $P_0$ and $P_1$ ) Synchronization: Algorithm 1

```

int turn:=0; /* j=1 - i, 0  $\leq$  i, j  $\leq$  1

 $P_i$  : repeat
    while (turn != i) do no-op;
    CS
    turn := j;
    RS
until false;

```

```

int turn:=0;

 $P_j$  : repeat
    while (turn != j) do no-op;
    CS
    turn := i;
    RS
until false;

```

4

## Two-Process Synchronization: Algorithm 2

```
boolean flag[2]; /* shared variables; initially flag[0] = flag[1] = false
                /* flag[i] == true  $\Rightarrow$  Pi can enter its critical section
```

**P<sub>i</sub> : repeat**

```
    while ( flag[j] ) do no-op;
    flag[i] := True;
```

critical section

```
    flag[i] := False;
```

remainder section

**until False;**

Correct? Satisfaction of (1) Mutual Exclusion? (2) Progress?  
(3) Bounded waiting?

5

## Two-Process Synchronization: Algorithm 2

```
boolean flag[2]:=False;
/* flag[i] == true  $\Rightarrow$ 
/* Pi can enter its CS.
```

**P<sub>i</sub> : repeat**

```
    while ( flag [ j ] ) do no-op;
```

```
    flag[i] := True;
```

CS

```
    flag[i] := False;
```

RS

**until False;**

```
boolean flag[2]:=False;
/* flag[j] == true  $\Rightarrow$ 
/* Pj can enter its CS.
```

**P<sub>j</sub> : repeat**

```
    while ( flag [ i ] ) do no-op;
```

```
    flag[j] := True;
```

CS

```
    flag[j] := False;
```

RS

**until False;**

6

## Two-Process Synchronization: Algorithm 3

```
boolean flag[2]; /* shared variables, initially flag[0] = flag[1] = false
                /* flag[i] == true  $\Rightarrow$  Pi ready to enter its critical section.
```

P<sub>i</sub> : repeat

```
    flag[i] := True;
    while ( flag[j] ) do no-op;
```

critical section

```
    flag[i] := False;
```

remainder section

until False;

Correct? Satisfaction of (1) Mutual Exclusion? (2) Progress?  
(3) Bounded waiting?

7

## Two-Process Synchronization: Algorithm 3

```
boolean flag[2]:=False;
/* flag[i] == true  $\Rightarrow$ 
/* Pi ready to enter its CS.
```

P<sub>i</sub> : repeat

```
    flag[i] := True;
    while ( flag[ j ] ) do no-op;
```

CS

```
    flag[i] := False;
```

RS

until False;

```
boolean flag[2]:=False;
/* flag[j] == true  $\Rightarrow$ 
/* Pj ready to enter its CS.
```

P<sub>j</sub> : repeat

```
    flag[ j ] := True;
    while ( flag[ i ] ) do no-op;
```

CS

```
    flag[ j ] := False;
```

RS

until False;

8

## Two-Process Synchronization: Algorithm 4 (Peterson's Algorithm)

```

boolean flag[2]; /* shared variables, initially flag[0] = flag[1] = false
                  /* flag[i] == true  $\Rightarrow$  Pi ready to enter its critical section.

int turn;
Pi: repeat

    flag[i] := True;
    turn := j;
    while (flag[j] && turn == j) do no-op;

    critical section

    flag[i] := False;

    reminder section
until False;

```

### Informal Proofs:

Mutual exclusion  
Progress  
Bounded waiting

Very fragile solution!

Meets all three requirements; solves the critical-section problem.

9

## Two-Process Synchronization: Algorithm 4

```

boolean flag[2]:=False;
/* flag[i] == true  $\Rightarrow$ 
   Pi ready to enter its CS.
int turn;

Pi: repeat
    flag[i] := True;
    turn := j;
    while (flag[j] && turn == j)
        do no-op;

    CS
    flag[i] := False;
    RS
until False;

```

```

boolean flag[2]:=False;
/* flag[j] == true  $\Rightarrow$ 
   Pj ready to enter its CS.
int turn;

Pj: repeat
    flag[j] := True;
    turn := i;
    while (flag[i] && turn == i)
        do no-op;

    CS
    flag[j] := False;
    RS
until False;

```

10

## Peterson's algorithm

Idea: Three-step solution:

1. Declare intention to enter into the CS.
2. Let others enter into their CSs.
3. Enter the CS.

This is an "example reasoning" for many critical section algorithms.

Where is the fragility of algorithm 4?

11

## Another Algorithm: Dekker's (much more complicated)

```

Shared variables:  int turn;           // 0 or 1
                  boolean flag[2];    // initially False

Pi:  repeat
    flag[i] = True;                    // I would like to enter.
    while (flag[j]) {                  // Does she?
        if (turn == j) {               // yes, and it is her turn.
            flag[i] = False;           // So I will yield
            while (turn == j);          // waiting for her to leave.
            flag[i] = True;            // Finally! She is done.
        }
    }

    critical section

    turn = j;
    flag[i] = False;                   // I had my turn.
                                        // Done

    remainder section

until False;

```

Proof: Mutual exclusion?  
Progress? Bounded waiting?

12

## Review: So Far, Brute-Force Process Synchronization Algorithms--Review (no synchronization constructs)

- So far we have seen:  
Very simple concurrent program structure, and  
two-process synchronization algorithms.
  - Algorithms 1 to 4. Correct one: fourth algorithm.
- Next: the same simple concurrent program structure, and  
n-process synchronization algorithm.

13

## N-process Synchronization: Bakery Algorithm (L. Lamport; 1974)

**Program structure:** As in the two process structure, except that there are now  $n$  processes,  $P_i$ ,  $0 \leq i \leq (n-1)$ .

**Solution Logic:** Implement a bakery-ticket mechanism (sort of).  
Before entering its critical section, a process gets a (ticket) number.  
But, multiple processes *can* have the same ticket number!

Holder of the smallest number enters the critical section.

If two processes  $P_i$  and  $P_j$  receive the same number and  $i < j$ ,  
then  $P_i$  enters into its CS first; else  $P_j$  enters its CS first.

14

## N-process Synchronization: Bakery Algorithm (L. Lamport; 1974)

**P<sub>i</sub>:** repeat

```

choosing[i] := True;
number[i] := max(number[0], number[1], ..., number[n - 1]) + 1;
choosing[i] := False;
for (j = 0; j < n; j++) {
    while (choosing[j]) do no-op;
    while (number[j] != 0 && (number[j], j) < (number[i], i))
        do no-op;
}

```

critical section

```
number[i] := 0;
```

remainder section

until False;

15

## N-process Synchronization: Bakery Algorithm

Shared data:

**boolean** choosing[n];

**int** number[n];

Initially, all choosing values are *False*;

All number values are 0.

**Notation:** lexicographical order (ticket #, process id #)

$(a, b) < (c, d)$  if  $a < c$  or  $(a = c \text{ and } b < d)$

16



## Bakery Algorithm: Proof of Mutual Exclusion

■ **Observation:** If

- (a)  $P_i$  is already in critical session, and
- (b) any  $P_k$ ,  $k \neq i$ , has chosen its number (i.e.,  $\text{number}[k] \neq 0$ ),  
then  $(\text{number}[i], i) < (\text{number}[k], k)$

Proof:

- Case 1:  $P_k$  **hasn't started to choose**  $\text{number}[k]$  when  $P_i$  executes  $\text{while}(\text{choosing}[k])$ . Obviously,  $\text{number}[i] < \text{number}[k]$ .
  - Case 2:  $P_k$  **has chosen**  $\text{number}[k] \neq 0$  before  $P_i$  executes  $\text{"while}(\text{choosing}[k])\text{"}$ .  
 $P_i$  wins, implying  $(\text{number}[i], i) < (\text{number}[k], k)$
  - Case 3:  $P_k$  **is choosing**  $\text{number}[k]$  when  $P_i$  executes  $\text{while}(\text{choosing}[k])$ .  
 $P_i$  will wait. We will reduce this to Case 2.
- Because of this observation, if  $P_i$  is already in its CS, no other process can enter into its CS.

17

## Why Do We Need choosing[ ]?

$P_i$ : repeat

```

choosing[i] := True;
number[i] := max(number[0], number[1], ..., number[n-1]) + 1;
choosing[i] := False;
for (j = 0; j < n; j++) {
    while (choosing[j]) do no-op;
    while (number[j] != 0 && (number[j], j) < (number[i], i))
        do no-op ;
}

```

critical section

```
number[i] := 0;
```

remainder section

until False;

18

## Why Do We Need choosing[ ]?

- Two processes, initially  $\text{number}[0]=\text{number}[1]=0$ .

$P_0$ :

evaluate  $\text{max}()$  function;  
... leave CPU...

..start executing...

$\text{number}[0] := 1$ ;  
while condition is false.  
enter critical section → violation of ME.

$P_1$ :

evaluate  $\text{max}()$  function;  
 $\text{number}[1] := 1$ ;  
enter critical section  
...leave CPU while in CS...

19

## Review: Brute-Force Process Synchronization Algorithms--Review (no synchronization constructs)

- So far, we have seen
  - Very simple program structure and two-process synchronization
    - ▶ Algorithms 1 to 4. Correct one: fourth algorithm.
  - Same simple program structure and n-process synchronization
    - ▶ Bakery Algorithm

20

## Hardware Solutions to Synchronization

- Up until now, we have seen brute-force ways of providing a solution to the critical section (CS) problem—only when the program structure is very limited.
- Next, we will see hardware-based solutions to the CS problem.

“Single CPU Cycle” (atomic) function as part of the hardware instruction set:

```
function Test-and-Set (var target: bool): bool;
begin
    Test-and-Set := target;
    target := True;
end;
```

21

## Mutual Exclusion with Test-and-Set

```
Shared data:
    bool lock := False;

Pi: repeat
    begin
        while (Test-and-Set (lock)) do no-op;

        critical section

        lock := False;

        remainder section
    end
until False;
```

Satisfaction of: Mutual exclusion? Bounded waiting? Progress?

22

## Another Hardware Sol'n: Swap

Also a “single CPU cycle” (atomic) procedure as part of the hardware instruction set:

```

Procedure Swap (var a, b: bool)
  var temp: bool;
  begin
    temp := a;
    a := b;
    b := temp;
  end;

```

23

## Mutual Exclusion with Swap

**Shared data:**

**bool** *lock* := *False*;

**P<sub>i</sub> : repeat**

**key** := *True*;

**repeat**

*Swap(lock, key)*;

**until** *key* = *False*;

/\* P<sub>j</sub> has its own key variable.

critical section

*lock* := *False*;

remainder section

**until** *False*;

Satisfaction of: mutual exclusion? Bounded waiting? Progress?

24

## Next: Three separate process synchronization constructs

- Semaphores
- Critical region or conditional critical region statements—CCRs  
--Only briefly. No CCR-related exam/assignment questions.
- Monitors
  - Btw, every Java object is/has a monitor: more on this later.
- NOTE: NO MIXING OF THESE CONSTRUCTS in your algorithms!
  - Your concurrent process algorithm can only use ONE of the three constructs.
    - ▶ If your algorithm uses semaphores, it cannot use CCRs or monitors.
    - ▶ If your algorithm uses CCRs, it cannot use semaphores or monitors.
- NOTE: ALWAYS USE THE SYNTAX SPECIFIED IN SLIDES!!

25

## Binary Semaphores

- Two *atomic operations* on a variable (semaphore) S:
- S below is a binary semaphore (takes on 0 or 1 as value):
 

$wait(S): \text{ while } S \leq 0 \text{ do no-op;}$ 
Dijkstra: P(S)  
 $S := S - 1;$

$signal(S): S := S + 1;$ 
Dijkstra: V(S)
- Semaphore S can **only** be accessed via **wait()** and **signal()**  
--two indivisible (i.e., atomic) operations.
- Initialization of S: either 0 or 1. Normally, 1.
- For a nice intro to semaphores, see the [semaphores chapter](#) in a [online book](#).

26

## Critical Section of $n$ Processes

Shared data:

**binary semaphore** mutex;    /\* mutex is set to 1 initially.

$P_i$ : repeat

`wait (mutex);`

critical section

`signal (mutex);`

remainder section

**Until** False;

Mutual exclusion?   Progress?   Bounded waiting?

27

## Problems with Binary Semaphores

**Binary semaphore S:**

`wait (S): while  $S \leq 0$  do no-op;`  
                    $S := S - 1$ ;

`signal (S):  $S := S + 1$ ;`

1. Busy waiting (i.e., no-op) above (in `wait(S)`) is wasteful.
2. Semaphore  $S$  above is binary.

Nonbinary (counting) semaphore operations form a **higher-level synchronization mechanism** that does not require busy waiting.

28

## Eliminating Busy Waiting and Bounded Waiting

**type** *nonbinary semaphore* = **record**

*value*: integer;                   /\* *value* is usually  
L: list of process;               /\* set to 1 initially  
**end**;

*wait*(S):

*S.value* := *S.value* - 1;

**if** (*S.value* < 0)

**then begin**

add this process to *S.L*;

*block*;

**end**;

*signal*(S):

*S.value* := *S.value* + 1;

**if** (*S.value* ≤ 0)

**then begin**

remove a process *P* from *S.L*;

*wakeup*(*P*);

**end**;

**Two simple operations:**

- *block* suspends the process that invokes it.
- *wakeup*(*P*) resumes the execution of a blocked process *P*.

29

## Producer-Consumer (Bounded Buffer) Problem

**Initialization:**

**binary semaphore** *mutex*:=1;

**nonbinary semaphore** *full*:=0; *empty*:=*n*;

**Producer:**

**repeat**

...

Produce an item in *nextp*;

*wait* (*empty*);   /\* When *empty* ≤ 0,

*wait* (*mutex*);   /\* block producer.

add *nextp* to buffer;

*signal* (*mutex*);

*signal* (*full*);

**until** *False*;

**Consumer:**

**repeat**

*wait*(*full*);   /\*When *full* ≤ 0; block consumer.

*wait* (*mutex*);

remove an item from buffer to *nextc*;

*signal*(*mutex*);

*signal*(*empty*);

...

consume the item in *nextc*;

**until** *False*;

**Proving Correctness:**

**Invariant I** (single-producer and single-consumer):

$$(n - 2) \leq \text{empty} + \text{full} \leq n \text{ and } 0 \leq \text{mutex} \leq 1$$

**Safety:** Prove that  $I \rightarrow \neg P_{\text{bad}}$

30

## Producer-Consumer (Bounded Buffer) Problem

Initialization:

binary semaphore  $mutex:=1$ ;

nonbinary semaphore  $full:=0$ ;  $empty:=n$ ;

**Producer:**

**repeat**

...

Produce an item in nextp;

$wait(empty)$ ; /\* When  $empty \leq 0$ ,

$wait(mutex)$ ; /\* block producer.

add nextp to buffer;

$signal(mutex)$ ;

$signal(full)$ ;

**until** False;

**Consumer:**

**repeat**

$wait(full)$ ; /\*When  $full \leq 0$ ; block consumer.

$wait(mutex)$ ;

remove an item from buffer to nextc;

$signal(mutex)$ ;

$signal(empty)$ ;

....

consume the item in nextc;

**until** False;

Note the symmetrical use of empty and full semaphores in entry and exit sections of producer and consumer.

31

## Dining Philosophers Problem (1965)

- 5 philosophers; 5 chopsticks; one bowl of rice.

nonbinary semaphore  $chopstick[0..4]:=1$ ;

/\* All chopsticks are initialized to 1.





## Dining Philosophers Problem

- A philosopher eats or thinks!
  - Each philosopher needs two chopsticks to eat, and chopsticks are picked up one at a time.
  - After successfully picking up two chopsticks, a philosopher eats for a while, and, then, puts down the chopsticks, and thinks.
- Problem: Decide how philosophers pick the chopsticks so that
  - Philosophers (processes) do not starve!
  - There is no deadlock.

## Dining Philosophers Problem

**nonbinary** *chopstick*[0..4]:=1;

**Philosopher *i*:**

**repeat**

*wait* (*chopstick*[*i*]);  
*wait* (*chopstick*[(*i*+1) mod 5]);

...

**EAT**

...

*signal* (*chopstick*[*i*]);  
*signal* (*chopstick*[(*i*+1) mod 5]);

...

**THINK**

...

**until** *False*;

**Problem: Deadlock**

## Dining Philosophers Problem

### Invariant property characterization:

up[i] = count of the times chopstick i has been picked up.

down [i] = count of the times chopstick i has been put down.

A chopstick cannot be put down more times than it has been picked up.

### Invariant property:

For all i:  $[ 1 \leq i \leq 5 \rightarrow (\text{down}[i] \leq \text{up}[i] \leq (\text{down}[i] + 1)) ]$

Note that a chopstick can be picked up by at most one philosopher at a time.

35

## Dining Philosophers Problem—No Deadlocks

**nonbinary semaphore** *chopstick*[ 0:4]:=1;

**Philosopher i:**

**repeat**

*wait (chopstick[(i+1) mod 5]);*  
*wait (chopstick[i]);*

...

**eat**

...

*signal (chopstick[i]);*  
*signal (chopstick[(i+1) mod 5]);*

...

**think**

...

**until** *False*;

Switched these two  
lines for philosopher 5,  
and NO deadlocks!

## To Avoid Deadlocks-One Approach

- A number of processes ( $P_i$ ) want to access an arbitrary number of resources ( $R_i$ ), needing exclusive access. Processes must abide by the following rules:
  - **Circular wait elimination via ordered resource requests:**  
Processes request resources in increasing order of resource ids. A philosopher obtains a low-numbered fork first.
  - **Eliminating no preemption:** If a process has obtained a resource  $R_i$  and needs resource  $R_j$ , it must release  $R_i$  before making the request if  $i > j$ .

## Readers-Writers Problem (1971)

**Assumption:** No reader waits unless a writer is writing! **Writers may starve!**

**Initialization:**

**nonbinary semaphore**  $mutex:=1$ ;  $wrt:=1$ ;

**int**  $readcount:=0$ ;

**Reader:**

```
wait (mutex);
  readcount := readcount + 1;
  if readcount = 1 then wait (wrt);
signal (mutex);
```

**In CS: Reading!**

```
wait (mutex);
  readcount := readcount - 1;
  if readcount = 0 then signal (wrt);
signal (mutex);
```

**Writer:**

```
wait (wrt);
  In CS: Writing!
signal (wrt);
```

## Readers-Writers Problem

**Invariant:** Use two auxiliary variables that characterize the state of concurrent execution.

writing [j] = 1 if writer j is in CS, 0 otherwise.

reading = 1 if **any** reader is in its CS, 0 otherwise.

Initially, reading=writing[j]=0 for all j,  $1 \leq j \leq n$

**Invariant I:** reading + writing [1] + ... + writing [n]  $\leq 1$

**Safety:** Prove that  $I \rightarrow \neg P_{\text{bad}}$

39

## Readers-Writers Problem-Nobody Starves.

**Assumption:** No reader waits unless a writer is writing!

Initialization:

nonbinary semaphore *mutex*:=1; *wrt*:=1; ***rw*:=1;**

int *readcount*:=0;

**Reader:**

```
wait(rw);
wait(mutex);
    readcount := readcount + 1;
    if readcount = 1 then wait (wrt);
signal (mutex);
signal(rw);
```

...

**In CS: Reading!**

```
...
wait (mutex);
    readcount := readcount - 1;
    if readcount = 0 then signal (wrt);
signal (mutex);
```

**Writer:**

```
wait(rw);
wait (wrt);
```

**In CS: Writing!**

```
signal (wrt);
signal(rw);
```

**Soln is NOT  
symmetric for  
readers and writers.  
Not good.**

40

## Yet Another Readers-Writers Problem: Readers Starve!

### Specifications:

- If a reader finds that there are writers waiting, the reader must wait until all waiting writers are finished writing.
- Even if writers arrive after the reader does, writers write before the reader reads, i.e., the reader can read only if the writers queue is empty.
- Readers can starve.

### Readers-Writers Problem

Give a semaphore-based solution to the Readers-Writers problem where “readers starve” in the sense that a stream of writers can arrive and “write” one after another even when a reader has been waiting to read.

#### Global variables:

semaphore mutex: initially set to 1    rmutex, wtr, rdr : all initially set to 0  
int nreaders = 0, nwriters = 0;    boolean Busy = False;    RBlocked = False;

#### Reader:

```
while True
{ wait (r-mutex);
  wait (mutex);
  if (nwriters > 0)
  then {RBlocked=True;
        signal(mutex);wait(rdr)}
  else {nreaders++;signal(mutex)}
};
signal(r-mutex);
READ;
wait(mutex);
nreaders--;
if (nreaders = 0 and
    nwriters > 0)
then {Busy = True; signal(wtr)};
signal(mutex);
DO-SOMETHING;
}
```

#### Writer:

```
while True
{ wait(mutex);
  nwriters++;
  if ( Busy or nreaders > 0)
  then {signal(mutex); wait(wtr)}
  else { Busy = True;
        signal(mutex)};
  WRITE;
  wait(mutex);
  nwriters--;
  Busy = False;
  if (nwriters > 0)
  then {Busy = True; signal(wtr)}
  else if (RBlocked) then
  {RBlocked = False;
   nreaders++; signal(rdr)};
  signal(mutex);
  DO-SOMETHING;
}
```

## Other Problems: Sleeping Barber Problem

- A barbershop consists of a waiting room with  $N$  chairs. If there are no customers to be served, the barber goes to sleep. If a customer enters the barbershop and all chairs are occupied, then the customer leaves the shop. If the barber is busy, but chairs are available, then the customer sits in one of the free chairs. If the barber is asleep, the customer wakes up the barber. Write a program to coordinate the barber and the customers.
- Objectives
  - No starvation: a customer will have his hair cut once he sits down.
  - No deadlocks: both barber and customers could be waiting.

43

## Cigarette-Smokers Problem

Consider a system with three smoker processes and one agent process. Each smoker continuously rolls a cigarette and then smokes it. But to roll and smoke a cigarette, the smoker needs three ingredients: tobacco, paper, and matches. One of the smoker processes has paper, another has tobacco, and the third has matches. The agent has an infinite supply of all three materials. The agent places two of the ingredients on the table. The smoker who has the remaining ingredient then makes and smokes a cigarette, signaling the agent on completion. The agent then puts out another two of the three ingredients, and the cycle repeats. Write a program to synchronize the agent and the smokers.

44

## Exercises

- Three processes,  $P_0$ ,  $P_1$ , and  $P_2$ . Write code with semaphores to synchronize them such that  $P_2$  executes only after both  $P_0$  and  $P_1$  execute.
- Two processes,  $P_0$  and  $P_1$ . Write code with semaphores to synchronize them such that  $P_0$  executes at least twice as often as  $P_1$  executes.